



Arpit Raj <arpitimp18@gmail.com>

Subject: Clarification on LLM use — Problem Statement 1 (Prototype Phase)

5 messages

Arpit Raj <arpitimp18@gmail.com>
To: support@hackerearth.com

17 June 2026 at 22:01

Hi Team,

We're a Prototype Phase (Round 2) team working on Problem Statement 1 — Poor Visibility on Parking-Induced Congestion. Before we finalize our submission, we'd like to confirm one point so we stay fully within the rules.

We understand from the FAQ that for Problem Statements 1 and 2 we must use only the dataset provided by HackerEarth, and that using external datasets may lead to disqualification. We want to be clear that all of our analysis — hotspot detection, congestion-impact scoring, root-cause classification, and enforcement prioritization — is computed entirely from the provided violation dataset. No external dataset is used anywhere in our pipeline.

Separately, we plan to add a natural-language assistant feature using a hosted large language model (LLM) API, and we want to describe exactly how it is used:

- The LLM is used only as a reasoning and phrasing layer. It takes an officer's plain-language question (e.g., "where should I focus near Koramangala?"), converts it into a query against tables we have already computed from the provided dataset, and then phrases the retrieved results in natural language.
- It does not ingest, train on, or augment our solution with any external dataset. Every figure it presents traces back to the provided dataset.
- We treat the LLM strictly as a tool/model (comparable to a library such as scikit-learn), not as a data source.

Could you please confirm whether this use of an LLM API is permitted for Problem Statement 1, or whether it would be treated as "external data" under the rules? We'd like to be certain before submitting.

Thank you for the clarification.

Best regards,
Arpit Raj
arpitimp18@gmail.com

HackerEarth Inc <support@hackerearth.com>
Reply-To: HackerEarth Inc <support@hackerearth.com>
To: arpitimp18@gmail.com

17 June 2026 at 22:14

Hello Arpit,

Thank you for reaching out to us.

We appreciate the detailed explanation of your proposed approach for Problem Statement 1 – Poor Visibility on Parking-Induced Congestion.

As per the guidelines shared for the Prototype Phase, participants working on Problem Statements 1 and 2 are required to use only the datasets provided as part of the challenge. The use of external datasets may result in disqualification and the submission will not be considered for evaluation.

The final review and evaluation of submissions will be carried out by the jury based on the challenge guidelines and submission requirements.

For reference, participants must use only the datasets provided for Problem Statements 1 and 2 and should avoid incorporating external datasets into their solution.

We hope this clarifies the currently published guidelines.

Regards,
Shruti
Technical Support Associate



[Quoted text hidden]

[Quoted text hidden]

To unsubscribe from this group and stop receiving emails from it, send an email to support+unsubscribe@hackerearth.com.

Arpit Raj <arpitimp18@gmail.com>
To: HackerEarth Inc <support@hackerearth.com>

18 June 2026 at 05:02

Hi Shruti / Support Team,

Thank you for your earlier response. I appreciate you taking the time to reply. However, I'd like to respectfully flag that my original question may not have been fully addressed. My concern was not about using an external dataset — I explicitly confirmed we are not doing that. My question was specifically about whether using a hosted LLM API (e.g., OpenAI, Anthropic) as a reasoning/phrasing layer constitutes an “external dataset” under the rules.

To make the distinction as clear as possible:

- Our entire analysis — hotspot detection, congestion scoring, root-cause classification — is computed solely from the HackerEarth-provided dataset.
- The LLM API is used only to interpret an officer's plain-language question and rephrase already-computed results in natural language.
- The LLM does not introduce, ingest, or draw from any external data source. It acts purely as a language interface over our pre-computed tables.

The core question is: Does using an LLM API as a reasoning/UI layer — with no external data involved — violate the “no external datasets” rule?

A yes/no answer with a brief rationale would be very helpful so we can finalize our submission with confidence.

Thank you for the clarification.

Best regards,
Arpit Raj
arpitimp18@gmail.com

[Quoted text hidden]

HackerEarth Inc <support@hackerearth.com>
Reply-To: HackerEarth Inc <support@hackerearth.com>
To: Arpit Raj <arpitimp18@gmail.com>

18 June 2026 at 05:58

Hello,

Thank you for reaching out to us!

We have received your request and would like to acknowledge it. We will get back to you shortly with an update.

Regards,
Vamsi,
Technical Support Associate



[Quoted text hidden]

HackerEarth Inc <support@hackerearth.com>
Reply-To: HackerEarth Inc <support@hackerearth.com>
To: Arpit Raj <arpitimp18@gmail.com>

18 June 2026 at 07:41

Hello Arpit,

Thank you for your patience!

Yes, participants may use LLM APIs as a reasoning/UI layer, provided they are used only to interpret user queries and generate natural-language responses based on outputs derived from the HackerEarth-provided dataset. This is permitted as long as no external datasets, knowledge sources, or retrieval mechanisms are used, and all analysis and insights are generated solely from the provided dataset. We hope this clarifies your concern.

If you have any other concerns, please get back to us.

Regards,
Vamsi,
Technical Support Associate



[Quoted text hidden]